



Kubernetes from 0 to production in 45 minutes

Bastian Hofmann

Agenda

- 01 • Simple web app
- 02 • Deploy it to Kubernetes
- 03 • You will learn about how
Kubernetes works along the way

What I do

What is Qdrant?

Qdrant is a vector similarity search engine
(or vector database)

- ◆ Semantic search
- ◆ Recommendations
- ◆ Fraud detection
- ◆ Anomaly detection
- ◆ Generative AI
- ◆ ...





kubernetes

Common compute platform across infrastructure

- ✓ Common API & Packaging
- ✓ Health Checks/HA
- ✓ Load Balancing
- ✓ Overlay Networking
- ✓ Network Security Policies
- ✓ Backup and Recovery
- ✓ Autoscaling
- ✓ Service Discovery
- ✓ Networking
- ✓ RBAC & Access Control




About containers

```
1 package main
2
3 import (
4     "fmt"
5     "net/http"
6 )
7
8 func hello(w http.ResponseWriter, req *http.Request) {
9     fmt.Fprintf(w, "Hello, world!\n")
10 }
11
12 func headers(w http.ResponseWriter, req *http.Request) {
13     for name, headers := range req.Header {
14         for _, h := range headers {
15             fmt.Fprintf(w, "%v: %v\n", name, h)
16         }
17     }
18 }
19
20 func main() {
21     http.HandleFunc("/hello", hello)
22     http.HandleFunc("/headers", headers)
23     http.ListenAndServe(":80", nil)
24 }
```

Application Code



Container Image



Application Dependencies


Demo

Containers are great



Managing a couple – no problem

Containers are great but



How about managing many?

How do we address:

Networking, Security, Scheduling, Automation, etc?

Container Orchestration

Kubernetes Architecture



Declarative API

Everything is a RESTful resource

CRUD on those resources

Pod

- One instance of your application
- Contains one or more containers
- Has an individual IP address within the Kubernetes network



Pod

```
apiVersion: v1
kind: Pod
metadata:
  name: my-application
  namespace: application-namespace
  labels:
    app: my-application
spec:
  containers:
    image: company/my-image:1.0.0
    ports:
      - containerPort: 8080
status:
  ...
```



Deployment

- Manages lifecycle of Pods
- Multiple replicas
- Different update strategies



Deployment



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-application
  namespace: application-namespace
spec:
  replicas: 2
  strategy:
    type: RollingUpdate
  template:
    spec:
      containers:
        - image: company/my-image:1.0.0
          ports:
            - containerPort: 8080
```

Service

- Kubernetes internal “load balancer”
- Static internal IP address mapped to healthy Pods
- Has a DNS entry in the Kubernetes internal DNS
- Different types to also expose workloads outside of the cluster



Service

```
apiVersion: v1
kind: Service
metadata:
  name: my-application
  namespace: application-namespace
spec:
  ports:
    - port: 80
      targetPort: 8080
  selector:
    app: my-application
  type: ClusterIP
```



Ingress

- Configuration of an Ingress Controller to route incoming traffic to the cluster to internal Services based on HTTP host and path
- Configure TLS



Ingress

```
apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: my-application
  namespace: application-namespace
spec:
  rules:
    - host: my-application.example.com
      http:
        paths:
          - backend:
              service:
                name: my-application
                port:
                  number: 8080
            path: /
            pathType: Prefix
```



ConfigMap & Secret

- Internal key/value store for configuration settings
- Can be used in Pods through environment variables or volume mounts




ConfigMap & Secret

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: my-application
  namespace: application-namespace
data:
  CONFIG_KEY: some-value
  ANOTHER_KEY: another value
```




PersistentVolumeClaim & PersistentVolume

- Mount persistent storage into Pods



PersistentVolumeClaim & PersistentVolume

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-application-volume
  namespace: application-namespace
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 8Gi
  storageClassName: some-storage-provider
```



There are many more

CustomResourceDefinition

- Extend the Kubernetes API with additional resource types
- Kubernetes operators can then implement additional abstractions and logic



Demo

Resources

- <https://k3s.io/>
- <https://k3d.io/>
- <https://longhorn.io/>
- <https://www.rancher.com/>
- <https://cert-manager.io/>
- <https://artifacthub.io/packages/helm/bitnami/redis>
- <https://helm.sh/>
- <https://kubernetes.io/>



Thank you

Bastian.Hofmann@qdrant.com